

Paper 6

AI-Assisted Portfolio Triage for Canada's Major Projects: Deploying a Risk Engine Dashboard

Running Title: AI Portfolio Triage – Dashboard Deployment

Author: Joshua Samuel

Affiliation: Northeast Renewables LP, Toronto, Ontario, Canada

Contact Information: Email: jsamuel@northeastmidstream.com

Keywords: Major Projects Inventory; Bayesian scorecard; Cox proportional hazards; Power ranking; MPI Dashboard; Project triage

Word Count: ~2,400

Date: August 2025

Funding Statement: This research received no external funding.

Conflict of Interest: The author declares no conflict of interest.

Abstract

Canada's Major Projects Inventory (MPI) records more than 500 large-scale energy, mining, and forestry proposals but provides little forward visibility on which will advance. Previous work established the descriptive and historical baseline and introduced two probabilistic models, a Bayesian scorecard, a Cox proportional-hazards model, as well as a composite power ranking. This paper documents the deployment of a minimum viable product (MVP): the MPI Dashboard. The dashboard ingests MPI data, applies pre-computed probabilities and urgency measures, and presents blended probabilities, priority indices, and power rankings through five interactive tabs: Probability and Ranking, Sector and Cleantech, Start-Year and Cost, Map, and Stage Flow. Performance testing showed sub-second refresh times and reproducible outputs. The MVP demonstrates that probabilistic triage can be transparent, responsive, and scalable, bridging the upstream intelligence gap and transforming a static registry into a living tool for planning and resource allocation.

Introduction

Canada's major project pipeline now exceeds \$1.3 trillion in announced value, yet historical data show that on average one in four proposals move from planning to construction within five years. This gap between ambition and execution risks stranded capital, stakeholder fatigue, and missed climate targets. Existing tools for project oversight, such as Project Portfolio Management (PPM) systems and Governance, Risk, and Compliance (GRC) platforms, typically activate only after a project has passed due diligence or received formal onboarding. This leaves a critical upstream intelligence gap: before an investment decision, sponsors and policymakers lack transparent, repeatable methods to triage large numbers of proposals quickly and defensibly.

Papers 1 and 2 established the descriptive and historical foundation of the MPI, documenting transition pathways, cohort survival, and bottlenecks in Canada's major project pipeline. Papers 3 through 5 then developed the statistical foundations of a forward-looking framework. Paper 3 introduced two lightweight, auditable models: a Bayesian likelihood-ratio scorecard, which estimates the probability that a project will reach construction, and a Cox proportional-hazards model, which estimates the time remaining to that milestone. Paper 4 operationalised these models at the portfolio level, generating blended probabilities and urgency-adjusted priority indices for hundreds of projects in the 2024 MPI. Paper 5 extended this work by introducing power ranking as a composite score that integrates probability and urgency into a single metric for resource allocation across competitive envelopes, demonstrated through a Clean Fuels Fund use case.

This sixth paper advances the framework from research to deployment. It documents the implementation of an MPI Dashboard, a minimum viable product (MVP) platform. The

dashboard integrates the Bayes and Cox risk engines, blended probability, priority index, and power ranking into an interactive interface that allows users to filter, rank, and visualize pre-construction projects. By providing real-time, data-driven triage, the MPI Dashboard transforms a static registry into an operational cockpit.

The contribution is twofold. First, it demonstrates how probabilistic risk models can be embedded in a production-grade dashboard that runs quickly, refreshes automatically, and remains fully auditable. Second, it shows how such a system overcomes the upstream intelligence gap, giving developers, planners, financiers, and policymakers a transparent tool for sequencing projects, targeting scarce capital, and aligning policy attention with the projects most likely and most urgent to advance.

Methods

Design basis

The MPI Dashboard was developed as a minimum viable product to demonstrate how probabilistic risk models could be embedded into a transparent, user-facing tool. Its design was guided by four principles.

Integration: The dashboard needed to combine the outputs of the Bayesian likelihood-ratio scorecard and the Cox proportional-hazards model into blended probabilities, urgency measures, and power rankings that could be visualized interactively.

Accessibility: Users required a low-friction interface, accessible through a browser without coding, with filters for province, sector, sponsor, cleantech status, cost band, and year.

Responsiveness: Performance targets were set at under two seconds for initial load and under one second for chart refreshes, even with typical portfolio sizes of 350–400 projects.

Auditability: Outputs had to remain reproducible, with fixed coefficient files, deterministic preprocessing rules, and CSV export for external validation.

System architecture

The system follows a modular pipeline with three layers.

Data ingestion and preprocessing: The dashboard ingests the MPI dataset in Excel format.

Required fields include categorical descriptors (province, sector, group, cleantech flag, status codes), project attributes (cost, start and end year, sponsor, project name, coordinates), and model outputs (blended probability, priority index, and power ranking). Input validation enforces schema consistency, coerces types, and preserves unknown values.

Risk engines: Probabilities are pre-computed externally by the Bayesian and Cox models, blended at a 60–40 ratio, and combined with years remaining to yield the priority index. Power ranking integrates blended probability and rescaled priority index into a single score.

Visualization layer: Built in Python using Plotly Dash, the interface exposes the dataset through a web application with dynamic filters, interactive charts, and downloadable outputs. Filters allow users to slice projects by province, sector, sponsor, cleantech status, cost band, and year. KPI cards at the top of the interface display the total number of active projects, the sum of investment value, and the average probability of construction. All views update automatically as filters change.

Results and Discussion

User Interface

The dashboard interface is organized into five main tabs, each offering a distinct analytical perspective on the dataset. The first tab, Probability & Priority, presents a scatter plot that positions projects according to their probability of construction and their urgency (Priority Index, normalized on a 0–1 scale). Median reference lines divide the chart into quadrants, allowing users to quickly distinguish between high and low probability or priority projects. Interactive hover functions display company, project, sector, province, estimated cost, and all three risk engine metrics—probability, priority, and power ranking—thereby providing detailed context for each observation.

The second tab, Power Ranking, offers a focused view on relative project rankings. It presents three Top-N horizontal bar charts that list projects by power ranking, probability, and priority index. Users can adjust the number of projects displayed through a slider, enabling targeted exploration of the most critical projects by each measure. This tab complements the scatter plot by providing explicit ranked lists that can be exported or referenced in downstream analysis.

The third tab, Facts & Figures, consolidates portfolio statistics in a two-by-four grid of charts. The first row shows stacked bar charts of projects by sector (broken down by subsector group) and by province (broken down by sector). The second row presents total project value by sector and by province. The third row contains two donut charts summarizing the share of projects by cleantech status (yes/no) and by ownership (public/private). The fourth row depicts the “vintage” of the portfolio, showing project counts or values by start year, stacked by sector. A Count vs Cost toggle in the sidebar controls whether Rows 1, 3, and 4 display project counts or aggregate

costs, while Row 2 is fixed to total cost. This tab provides both compositional and value-based perspectives on the portfolio.

The fourth tab, Map, situates projects geographically on an interactive base map using latitude and longitude coordinates. Each project is represented by a bubble whose size corresponds to its capital cost, while color encodes sector. This view emphasizes the regional distribution of projects and highlights geographic clustering of investment.

The fifth and final tab, Stage Flow, visualizes transitions between start and end status codes using a Sankey diagram. This representation illustrates how projects progress, stall, or attrit along the development pipeline. It highlights bottlenecks and common exit points, offering an intuitive view of project dynamics over time.

A persistent sidebar anchors the interface and provides a comprehensive suite of filters and controls that apply across all tabs. Users can filter projects by company, province, sector, subsector group, cleantech status, and project status, as well as select specific companies or projects by name. Additional sliders allow dynamic adjustment of year ranges and cost ranges, while the Count vs Cost toggle switches several charts between frequency and value-based displays. The sidebar also includes a Top-N slider for ranking charts, reset and download buttons, and contextual information about the loaded dataset. This common filtering logic ensures that all visualizations remain synchronized, enabling a coherent exploration of the dataset from multiple perspectives.

Multi-Panel Dashboard

Figure 1 presents the five-tab dashboard, illustrating the range of analytical perspectives available in a single interface. Tab 1 shows the Probability & Priority scatter plot with quadrant overlays. Tab 2 displays the Top-N bar charts by power ranking, probability, and priority index. Tab 3 illustrates the Facts & Figures grid, which combines sectoral, provincial, cleantech, ownership, and vintage breakdowns into a unified set of visuals. Tab 4 provides the geospatial map with proportional bubbles representing project cost, while Tab 5 shows the Stage Flow diagram tracing transitions between start and end status codes. These tabs taken together demonstrate the dashboard's ability to render multiple complementary views of the project pipeline, all linked by a common filtering logic.




Figure 1: Multi-Tab Dashboard.

Performance and reproducibility

Testing on MPI files of roughly 350–400 projects confirmed that the dashboard met its design basis. Initial load times averaged under two seconds, while filter and chart updates are refreshed in less than one second. CSV export produced identical outputs across sessions, ensuring reproducibility.

User insights

Interactive filtering enabled rapid triage across multiple dimensions. In the Probability and Ranking tab, quadrant overlays separated shovel-ready projects from long-horizon or low-likelihood files. Sector and cleantech charts revealed clustering patterns by region and industry. The geospatial tab highlighted provincial concentrations of high-priority projects, while the stage-flow diagram provided a systemic view of progression and attrition.

Lessons from the MVP

Deployment of the MPI Dashboard demonstrated that probabilistic risk models can be embedded into a responsive, user-facing system with relatively light infrastructure. Three lessons stand out.

Transparency is achievable at scale. By exposing blended probabilities, priority indices, and power rankings directly in the interface, the dashboard allows users to trace how each project's score is derived. This builds confidence among stakeholders in contrast to opaque black-box models.

Responsiveness matters as much as accuracy. Even simple models will not be adopted if users must wait for results. Meeting the sub-second refresh thresholds proved critical: planners and policymakers can experiment with filters in real time, encouraging exploration rather than static reporting.

Data quality remains the bottleneck. The risk engines are deterministic and reproducible, but the MPI dataset contains missing values, inconsistent labels, and limited contextual variables. The dashboard mitigates this by coercing unknown categories and enforcing schema rules, but future adoption will depend on cleaner, timelier data.

Practical constraints

The MVP revealed several constraints. Manual file upload complicates version control and makes it difficult to ensure that all users are working from the same dataset. In addition, rigid input schema requirements mean that project files must strictly match expected column names, which can create friction when integrating new or updated data. Contextual blind spots persist, as permitting milestones, Indigenous benefit agreements, or commercial contracts are not yet captured in the MPI dataset. These constraints restrict the current dashboard to exploratory use rather than operational integration.

Next steps

The path forward has three layers.

1. Data integration: Automating ingestion from official MPI releases would eliminate manual uploads and version drift. Adding contextual signals such as regulatory permits or commercial contracts would improve risk discrimination without overcomplicating the models.
2. Platform development: A lightweight API would allow users to score new projects and benchmark them against the existing dataset. Multi-user access, secure hosting, and user authentication would transition the dashboard from MVP to production-ready service.
3. Analytical extensions: Future builds could incorporate dynamic weighting between probability and urgency, simulate funding allocations across competitive envelopes, or link to external indicators. These enhancements would position the dashboard not just as a monitor but as a decision-support tool for governments, financiers, and developers.

Broader implications

The MPI Dashboard shows how upstream intelligence gaps can be bridged by embedding probabilistic models in transparent, interactive systems. It reduces analytic drag, lowers the cost of project surveillance, and provides a reproducible baseline for triage decisions. While further development is needed before operational deployment, the MVP confirms that a lean, auditable dashboard can transform a static project registry into a living tool for resource allocation and strategic planning.

Conclusion

This paper documented the implementation of the MPI Dashboard as a minimum viable product. It completes a sequence of work that began with Papers 1 and 2, which established the descriptive and historical foundations of the MPI through transition analysis, cohort tracking, and bottleneck identification. Papers 3 through 5 then developed the probabilistic risk framework: Paper 3 introduced the Bayesian scorecard and Cox proportional-hazards model, Paper 4 applied them at the portfolio level to generate blended probabilities and priority indices, and Paper 5 introduced power ranking as a composite score to order projects by both feasibility and urgency. Paper 6 brings these elements together in a functioning, user-facing system.

The dashboard met its design basis for integration, accessibility, responsiveness, and auditability. By combining Bayesian and Cox outputs into blended probabilities, priority indices, and power rankings, and presenting them through an interactive interface, it demonstrates how probabilistic intelligence can be delivered in real time. Performance testing confirmed that load and refresh times meet responsiveness thresholds, while reproducibility safeguards ensure consistent outputs across sessions.

Lessons from the minimum viable product are clear. Transparency is achievable at scale when scores are exposed directly to users. Responsiveness is critical for adoption, as real-time updates encourage exploration and trust. Data quality remains the limiting factor, requiring cleaner and timelier inputs to sustain credibility. The MVP also highlighted practical constraints, including manual uploads, rigid schema requirements, and contextual blind spots not yet captured in the MPI dataset.

Future development will focus on automating ingestion of official MPI releases, adding contextual signals such as permitting and commercial milestones, and enabling secure multi-user access through an API. Analytical extensions could allow dynamic weighting between probability and urgency, simulate funding allocations across competitive envelopes, or link to external indicators. With these enhancements, the dashboard can evolve from an exploratory tool into a decision-support cockpit.

By bridging the upstream intelligence gap that precedes conventional portfolio management and compliance systems, the MPI Dashboard transforms a static registry into a living analytical platform. Taken together, Papers 1 and 2 provided the baseline understanding, Papers 3 through 5 built and refined the risk models and ranking framework, and Paper 6 demonstrates deployment. The series establishes a transparent, auditable, and scalable framework for understanding Canada's major project pipeline and aligning capital and policy attention with the projects most likely, and most urgent, to advance.